

Chapter 6

Pipelining

Chapter Outline

- Pipelining: overlapped instruction execution
- Hazards that limit pipelined performance gain
- Hardware/software implications of pipelining
- Influence of pipelining on instruction sets

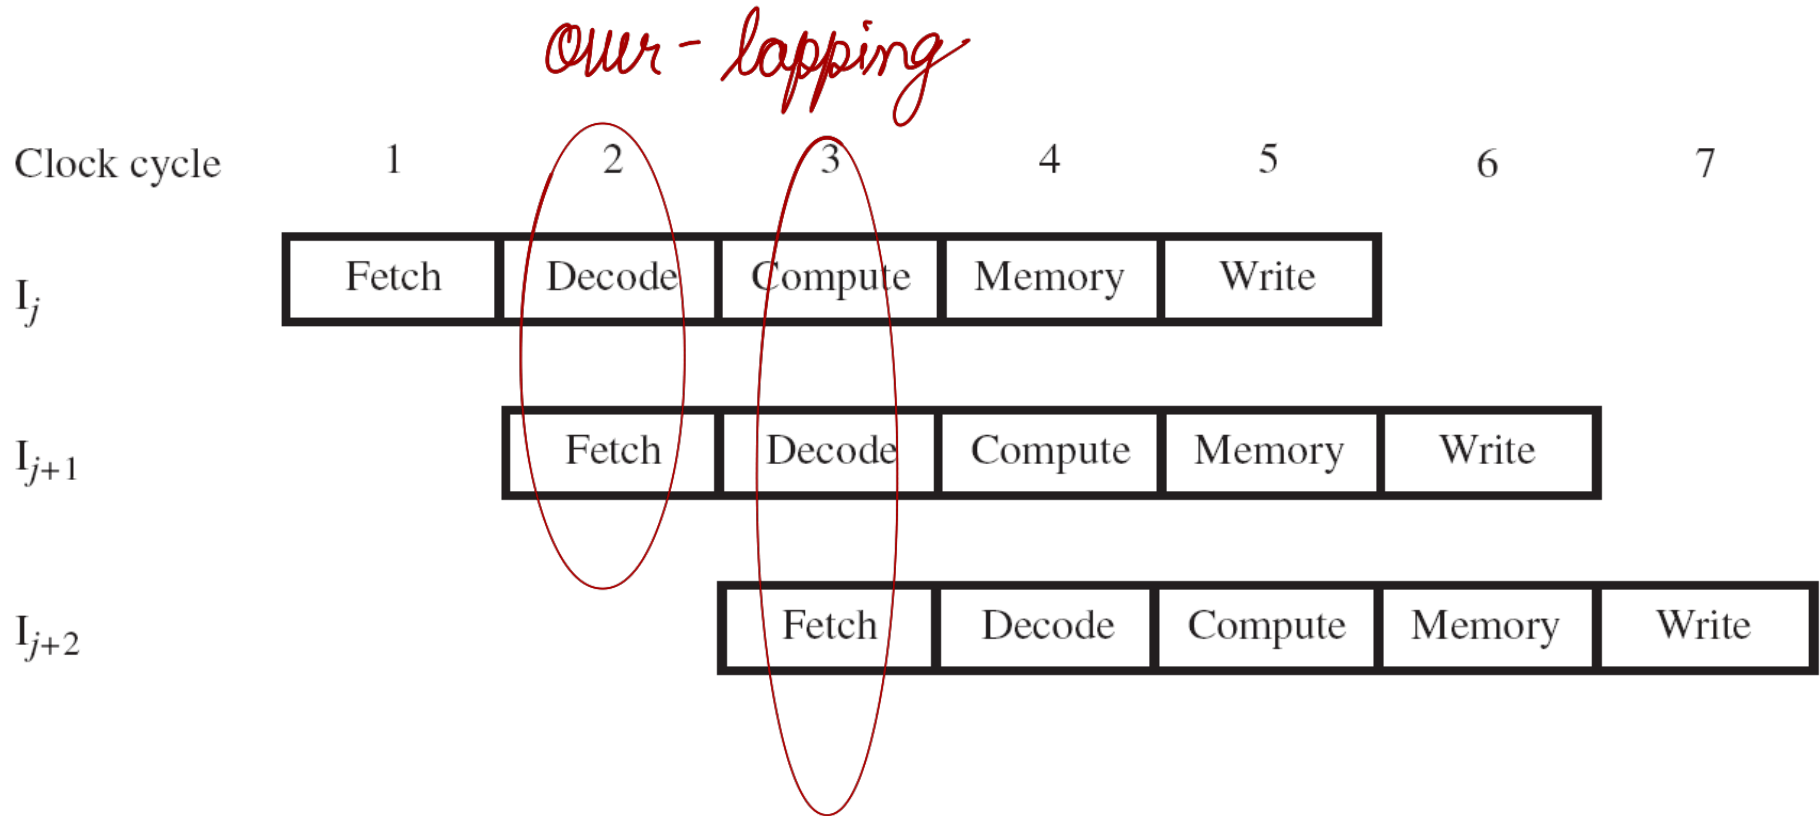
Basic Concept of Pipelining

- Circuit technology and hardware arrangement influence the speed of execution for programs
- All computer units benefit from faster circuits
- Pipelining involves arranging the hardware to *perform multiple operations simultaneously*
- Similar to assembly line where product moves through stations that perform specific tasks
- Same total time for each item, but *overlapped*

Pipelining in a Computer

- Focus on pipelining of *instruction execution*
- Multistage datapath consists of: Fetch, Decode, Compute, Memory and Write
- Instructions fetched & executed one at a time with only one stage active in any cycle
- *With pipelining*, multiple stages are active simultaneously for different instructions
- Still 5 cycles to execute, but *rate* is 1 per cycle

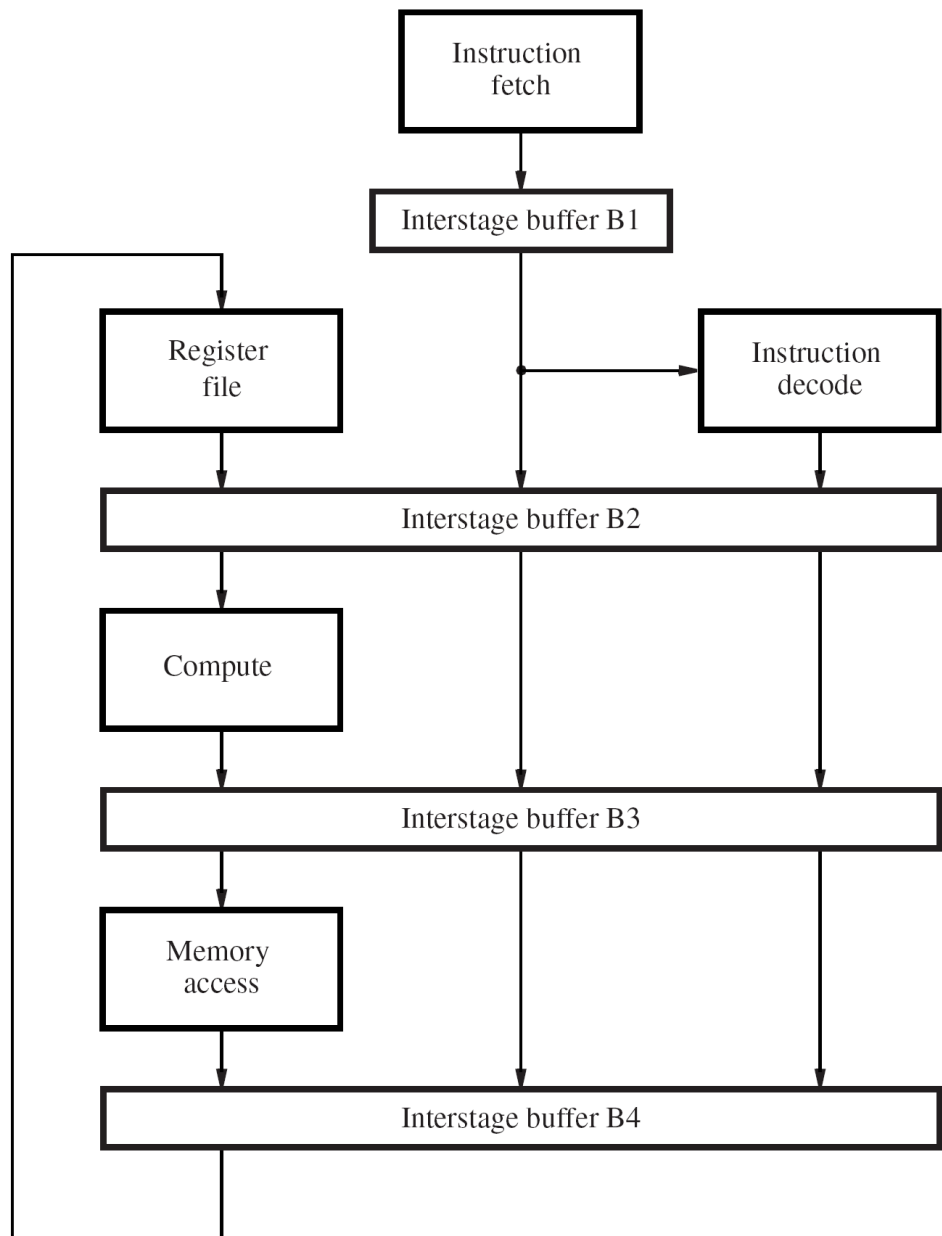
Pipelined Execution – Ideal case



At any given time, each stage of the pipeline is processing a different instruction.

Pipeline Organization

- Use program counter (PC) to fetch instructions
- A new instruction enters pipeline every cycle
- Carry along instruction-specific information as instructions flow through the different stages
- Use *interstage buffers* to hold this information
- These buffers incorporate RA, RB, RM, RY, RZ, IR, and PC-Temp registers discussed earlier
- The buffers also hold control signal settings



Datapath operands
and results

Source/destination
register identifiers
and other information

Control signals
for different stages

Pipelining Issues


- Consider two successive instructions I_j and I_{j+1}
- Assume that the destination register of I_j matches one of the source registers of I_{j+1}
- Result of I_j is written to destination in cycle 5
- But I_{j+1} reads *old* value of register in cycle 3
- Due to pipelining, I_{j+1} computation is incorrect
- So *stall* (delay) I_{j+1} until I_j writes the new value
- Condition requiring this stall is a *data hazard*

Data Dependencies

- Now consider the specific instructions
Add R2, R3, #100
Subtract R9, R2, #30
- Destination R2 of Add is a source for Subtract
- There is a *data dependency* between them because R2 carries data from Add to Subtract
- On *non*-pipelined datapath, result is available in R2 because Add completes before Subtract

Stalling the Pipeline

- With pipelined execution, old value is still in register R2 when Subtract is in Decode stage
- So **stall** Subtract for 3 cycles in Decode stage
- New value of R2 is then available in cycle 6



Details for Stalling the Pipeline


- Control circuitry must recognize dependency while Subtract is being decoded in cycle 3
- Interstage buffers carry register identifiers for source(s) and destination of instructions
- In cycle 3, compare destination identifier in Compute stage against source(s) in Decode
- R2 matches, so Subtract kept in Decode while Add allowed to continue normally

Details for Stalling the Pipeline

- Stall the Subtract instruction for 3 cycles by holding interstage buffer B1 contents steady
- But what happens after Add leaves Compute?
- Control signals are set in cycles 3 to 5 to create an *implicit* NOP (No-operation) in Compute
- NOP control signals in interstage buffer B2 create a cycle of idle time in each later stage
- The idle time from each NOP is called a *bubble*

Operand Forwarding

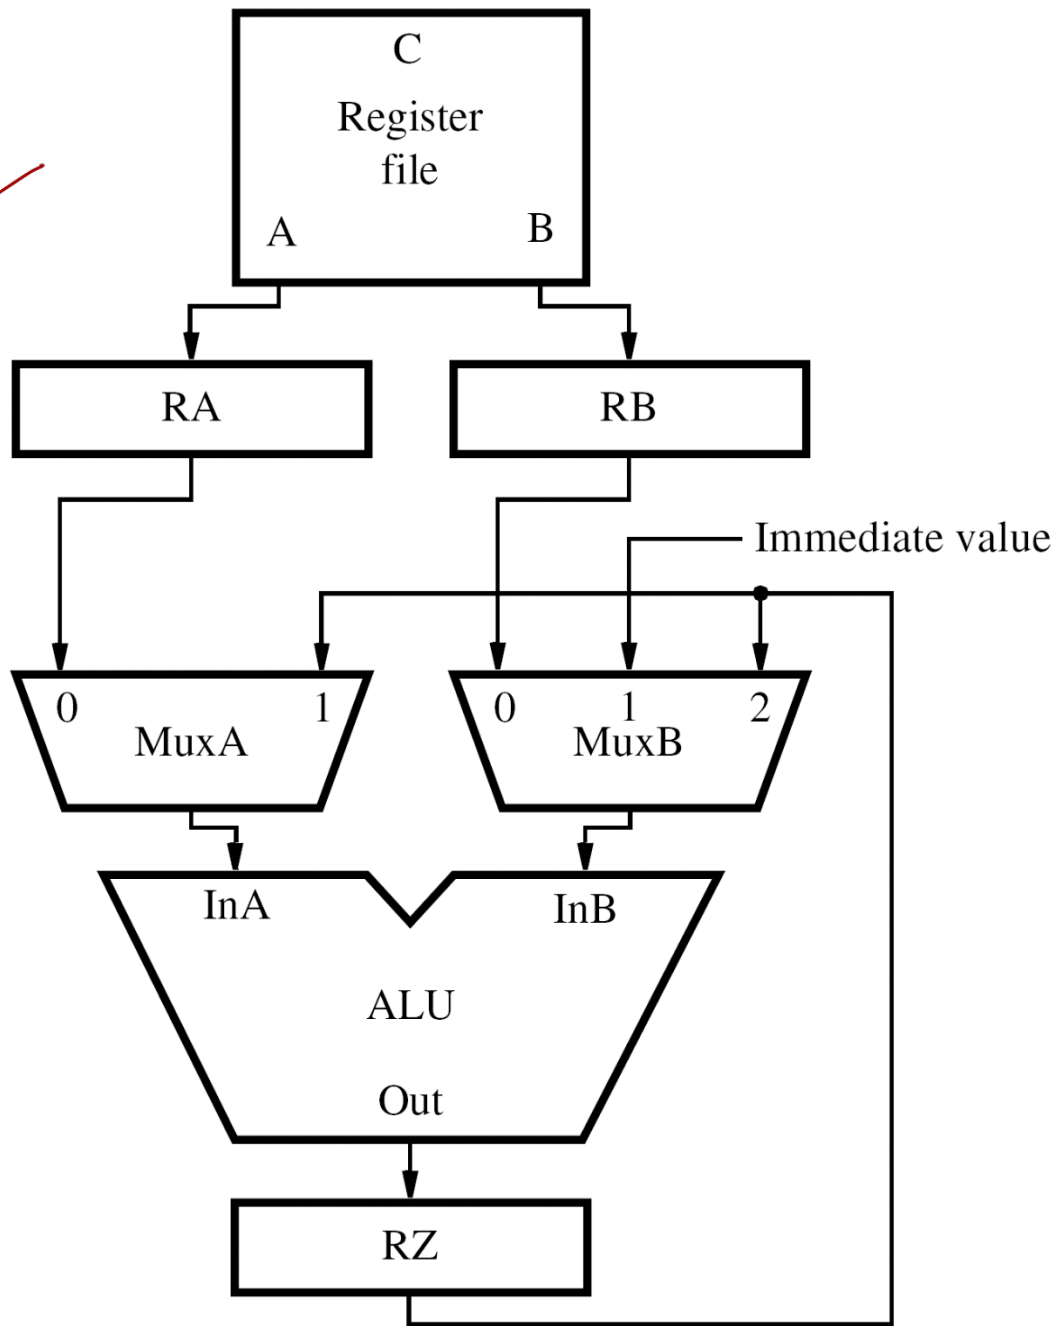

- **Operand forwarding** handles dependencies without the penalty of stalling the pipeline
- For the preceding sequence of instructions, new value for R2 is available at end of cycle 3
- *Forward* value to where it is needed in cycle 4





Details for Operand Forwarding

- Introduce multiplexers before ALU inputs to use contents of register RZ as forwarded value
- Control circuitry now recognizes dependency in cycle 4 when Subtract is in Compute stage
- Interstage buffers still carry register identifiers
- Compare destination of Add in Memory stage with source(s) of Subtract in Compute stage
- Set multiplexer control based on comparison



More on Operand Forwarding

Add	R2, R3, #100
Or	R4, R5, R6
Subtract	R9, R2, #30

- Consider another situation as above
- Introduce multiplexers before ALU inputs to use contents of register RY as forwarded value
- Extend MuxA and MuxB with this input as well


Software Handling of Dependencies

- Compiler can generate & analyze instructions
- Data dependencies are evident from registers
- Compiler puts three *explicit* NOP instructions between instructions having a dependency
- Delay ensures new value available in register but causes total execution time to increase
- Compiler can *optimize* by moving instructions into NOP slots (if data dependencies permit)

Software Handling of Dependencies

```
Add      R2, R3, #100
NOP
NOP
NOP
Subtract  R9, R2, #30
```


Insertion of NOP instructions for a data dependency



Memory Delays


- Memory delays can also cause pipeline stalls
- A cache memory holds instructions and data from the main memory, but is faster to access
- With a cache, typical access time is one cycle
- But a cache *miss* requires accessing slower main memory with a much longer delay
- In pipeline, memory delay for one instruction causes subsequent instructions to be delayed

Memory Delays



Memory Delays

- Even with a cache *hit*, a Load instruction may cause a short delay due to a data dependency
- One-cycle stall required for correct value to be forwarded to instruction needing that value
- Optimize with useful instruction to fill delay




Branch Delays – Control Hazard

- Ideal pipelining: fetch each new instruction while previous instruction is being decoded
- Branch instructions alter execution sequence, but they must be processed to know the effect
- Any delay for determining branch outcome leads to an increase in total execution time
- Techniques to mitigate this effect are desired
- Understand branch behavior to find solutions

Unconditional Branches

- Consider instructions I_j , I_{j+1} , I_{j+2} in sequence
- I_j is an unconditional branch with target I_k
- Compute stage normally determines target address using offset and PC+4 value
- In pipeline, target I_k is known for I_j in cycle 4, but instructions I_{j+1} , I_{j+2} fetched in cycles 2 & 3
- Target I_k should have followed I_j immediately, so discard I_{j+1} , I_{j+2} and incur two-cycle *penalty*


Unconditional Branches



Reducing Branch Penalty

- In pipeline, adder for PC is used every cycle, so it cannot calculate the branch target address
- So introduce a second adder just for branches
- Place this second adder in the Decode stage to enable earlier determination of target address
- For previous example, now only I_{j+1} is fetched
- Only one instruction needs to be discarded
- The branch penalty is reduced to one cycle

Reducing Branch Penalty



Conditional Branches

- Consider a conditional branch instruction:
 Branch_if_[R5]=[R6] LOOP
- Requires not only target address calculation, but also requires comparison for condition
- ALU normally performs the comparison
- Target address now calculated in Decode stage
- To maintain one-cycle penalty, we introduce a comparator just for branches in Decode stage

Branch Delay Slot

- Let both branch decision and target address be determined in Decode stage of pipeline
- Instruction immediately following a branch is always fetched, regardless of branch decision
- That next instruction is discarded with penalty, except when conditional branch is not taken
- The location immediately following the branch is called the *branch delay slot*

Branch Delay Slot

- Instead of conditionally discarding instruction in delay slot, *always* let it complete execution
- Let compiler find an instruction *before* branch to move into slot, if data dependencies permit
- Called *delayed branching* due to reordering
- If useful instruction put in slot, penalty is *zero*
- If not possible, insert explicit NOP in delay slot for one-cycle penalty, whether or not taken

Add	R7, R8, R9
Branch_if_[R3]=0	TARGET
I_{j+1}	
\vdots	
TARGET:	I_k

(a) Original sequence of instructions containing a conditional branch instruction

Branch_if_[R3]=0	TARGET
Add	R7, R8, R9
I_{j+1}	
\vdots	
TARGET:	I_k

(b) Placing the Add instruction in the branch delay slot where it is always executed

Branch Prediction


- A branch is decided in Decode stage (cycle 2) while following instruction is *always* fetched
- Following instruction may require discarding (or with delayed branching, it may be a NOP)
- Instead of discarding the *following* instruction, can we anticipate the *actual* next instruction?
- Two aims: (a) *predict* the branch decision
(b) use prediction *earlier* in cycle 1

Static Branch Prediction


- Simplest approach: assume branch *not* taken
- Penalty if prediction disproved during Decode
- If branches are random, accuracy is 50%
- But a branch at end of a loop is usually taken
- So for backward branch, always predict *taken*
- Use target address as soon as it is available
- Expect higher accuracy for this special case

Dynamic Branch Prediction

- Idea: track branch decisions during execution for *dynamic* prediction to improve accuracy
- Simplest approach: use most recent outcome for likely taken (LT) or likely not-taken (LNT)
- For branch at end of loop, we mispredict in last pass, and in first pass if loop is *re-entered*
- Avoid misprediction for loop re-entry with four states (ST, LT, LNT, SNT) for strongly/likely
- Must be wrong *twice* to change prediction



(a) A 2-state algorithm



(b) A 4-state algorithm

Resource Limitations

- Effect of Resource Limitation
 - Structural hazards
 - Both instruction fetch and operand fetch from memory at the same time
 - Use separate L1 cache for instruction and data

Performance Evaluation

- For non-pipelined processor, execution time T of a program with N instructions is given by

$$T = (N \times S)/R$$

S is the average number of clock cycles per instruction

R is the clock rate in cycles per second

- Throughput $P_{np} = R/S$
- For 5 stage RISC processor, with no pipelining, $S=5$
- For pipelined processor with no stalls, $S=1$ and hence, throughput $P_p = R$

Performance Evaluation

- Effect of Data Hazard
- δ_{stall} = Additional number of cycles per instruction due to stall over $S=1$
- Consider one-cycle stall due to data hazard (With Load instructions) with 25% of instructions being Load instructions and 40% of Load instructions followed by dependent instruction.
- $\delta_{\text{stall}} = 0.25 \times 0.40 \times 1 = 0.10$
- $P_p = R / (1 + \delta_{\text{stall}}) = R / 1.1 = 0.91R$

Performance Evaluation

- Effect of mis-predicting branch
- Consider one-cycle branch penalty with 20% of instructions being Branch instructions and average branch prediction accuracy is 90%.
- $\delta_{\text{branch_penalty}} = 0.20 \times 0.10 \times 1 = 0.02$
- Number of Pipeline Stages
- Large value increases throughput but also causes more hazards due to increased dependencies
- ALU is also pipelined for faster clock. Modern technology uses deep pipelines (More than 20)

Performance Evaluation

- Effect of cache miss
- Slower main memory stalls the pipeline for p_m cycles in the event of cache miss.
- Fraction m_i of all **instructions** that are fetched incur a cache miss.
- Fraction d of all instructions are Load or Store instructions, i.e., **data** to be loaded/stored.
- Fraction m_d of them incur cache miss.
- Increase over the ideal case of $S = 1$ due to cache misses is $\delta_{\text{miss}} = (m_i + d \times m_d) \times p_m$

Performance Evaluation

- Suppose that 5% of all fetched instructions incur a cache miss
- 30% of all instructions executed are Load or Store instructions, and 10% of their data-operand accesses incur a cache miss.
- Assume that the penalty to access the main memory for a cache miss is 10 cycles.
- Increase over the ideal case of $S = 1$ due to cache misses in this case is given by
- $\delta_{\text{miss}} = (0.05 + 0.30 \times 0.10) \times 10 = 0.8$

Performance Evaluation

- When all factors are combined, S is increased from the ideal value of 1 to

$$1 + \delta_{\text{stall}} + \delta_{\text{branch_penalty}} + \delta_{\text{miss}}$$

- Contribution of cache misses is often the dominant one.

Sections to Read

(From Hamacher's Book)

- Chapter on Pipelining
 - All sections and sub-sections upto and including Sub-section 6.8.2.
 - Not needed: Complete Sections 6.9 and 6.10